

Business, Economics, and Scaling

Monetization Model

NeuroTriage AI is designed as an **Enterprise Software-as-a-Service (SaaS)** solution for hospitals and diagnostic imaging centers. Hospitals subscribe annually based on the number of imaging sites or CT scanners covered. The subscription includes software updates, technical support, performance monitoring, and continuous model improvements. This model provides predictable operating costs for customers while enabling continuous product evolution.

Value Logic

NeuroTriage AI creates value by improving emergency workflow efficiency rather than replacing radiologists.

Stakeholder	Value Created
Patients	Earlier identification of suspected intracranial hemorrhage, enabling faster diagnosis and treatment.
Radiologists	Intelligent worklist prioritization, allowing critical studies to be reviewed earlier while reducing manual effort.
Hospitals	Improved emergency response, better utilization of radiology resources, and enhanced quality of patient care without disrupting existing workflows.
Product Provider	Long-term customer relationships and continuous product improvement through clinician feedback.

The product delivers value by enabling faster clinical decision-making while preserving clinician oversight and patient safety.

Cost Drivers

Cost Driver	Examples
Model / API Cost	Model development, inference infrastructure, API hosting, model updates
Support Cost	Hospital onboarding, user training, technical support, software maintenance
Data / Operations Cost	Data collection, anonymization, quality validation, monitoring, model retraining
Clinical Review Cost	Expert validation, pilot evaluation, Human-in-the-Loop review, regulatory documentation

Effective management of these costs is essential to ensure long-term product sustainability.

Unit Economics View

NeuroTriage AI requires significant upfront investment in AI development, clinical validation, and regulatory readiness. However, once deployed, the same AI platform can be scaled across multiple hospitals with relatively low incremental deployment cost.

Long-term sustainability depends on:

- High clinician adoption.
 - Measurable improvements in emergency workflow efficiency.
 - Continuous product improvement through real-world clinical feedback.
 - Efficient operational support across customer sites.
-

Scaling Risks

Risk	Potential Impact	Mitigation Strategy
Latency	Delayed worklist prioritization reduces clinician trust.	Optimize inference performance and infrastructure.
Cost Growth	Increasing infrastructure and maintenance costs as deployments expand.	Efficient resource utilization and scalable cloud architecture.

Workflow Complexity	Different hospitals use different PACS systems and workflows.	Standardized APIs and configurable product integration.
Governance Burden	Increased monitoring, compliance, and model oversight requirements.	Continuous performance monitoring and periodic governance reviews.

A phased deployment strategy is recommended, beginning with pilot implementation in a single hospital, followed by expansion to multiple hospitals before extending the platform to support additional neurological conditions.